

# Technical Work Report

LLM-Based HSE Safety Report Generation Module — Code Review, Schema Reconciliation & Demo Integration Guide

| Field        | Detail                                                                                    |
|--------------|-------------------------------------------------------------------------------------------|
| Subject      | Review of LLM.zip (llm_client.py, prompt.py, models.py, llm_main.py)                      |
| Prepared for | Demo Team                                                                                 |
| Date         | August 3, 2026                                                                            |
| Scope        | Full source review, gap analysis, and Pydantic schema alignment with pipeline JSON output |

## 1. Executive Summary

This report documents a full review of the four files delivered inside LLM.zip, which together form the LLM-based reporting layer of an HSE (Health, Safety & Environment) video-analytics pipeline. The module downloads and runs a local Qwen2.5-1.5B-Instruct language model on a GPU, and generates Markdown-formatted safety reports (daily, weekly, incident, and executive-summary variants) from structured detection data produced upstream by a video tracking / rule engine.

The review found the core generation logic (model loading and text generation) and the prompt templates to be sound and usable as-is. However, two issues require attention before the module is demo-ready:

llm\_main.py, the intended entry point that should tie model loading, data loading, and report generation together, is an empty file (0 bytes) — there is currently no runnable pipeline, only components.

The existing Pydantic schema in models.py (RuleEngineData / Event / Statistics) does not match the actual JSON structure produced by the tracking pipeline (video\_info / tracking\_summary / alerts\_report). This report defines a corrected schema, provided as models\_v2.py, that validates against the real output format.

Section 7 of this report provides the corrected Pydantic model, and Section 8 gives the Demo Team a concrete, step-by-step procedure for running the module end-to-end against a sample tracking JSON file.

## 2. Package Contents

LLM.zip contains the following four files, with no subfolders, README, or requirements file included:

| File          | Size    | Role                                                                         |
|---------------|---------|------------------------------------------------------------------------------|
| llm_client.py | ~8 KB   | Downloads and runs the local Qwen2.5 model (model loading + text generation) |
| prompt.py     | ~8 KB   | Builds the four report prompt templates sent to the LLM                      |
| models.py     | ~4 KB   | Pydantic input schema (currently out of date — see Section 6)                |
| llm_main.py   | 0 bytes | Empty — intended pipeline entry point, not implemented                       |

Note: no requirements.txt, setup.py, or \_\_init\_\_.py was found in the archive. The Demo Team will need to assemble these (see Section 8).

## 3. File-by-File Review

### 3.1 llm\_client.py

Provides two building blocks:

`download_model(path, model_name)` — downloads a Hugging Face model snapshot (default: Qwen/Qwen2.5-1.5B-Instruct) to a local folder via `huggingface_hub.snapshot_download`. Only needs to be run once per machine.

`LLMModel` — a wrapper class that loads the tokenizer and model from a local directory onto a CUDA GPU in bfloat16 precision, and exposes `generate(prompt, max_new_tokens) / __call__()` for single-turn chat-style generation using greedy decoding (`do_sample=False`).

Observations:

CUDA is a hard requirement — `__init__` raises `RuntimeError` immediately if no GPU is available. There is no CPU fallback path, so the demo machine must have an NVIDIA GPU with a CUDA-compatible PyTorch install.

Generation is single-turn only (one user message per call); there is no conversation/history support, which is appropriate for this report-generation use case.

Decoding is deterministic (`do_sample=False`), which is the right choice for a compliance/safety report — outputs are reproducible for the same input data.

### 3.2 prompt.py

Defines `HSEPromptGenerator`, a class of four `@classmethod` prompt builders, each taking a `RuleEngineData` object and returning a prompt string:

| Method                           | Output sections                                                       | Purpose                                                                           |
|----------------------------------|-----------------------------------------------------------------------|-----------------------------------------------------------------------------------|
| <code>daily_report()</code>      | Summary, Statistics, Key Violations, Risk Assessment, Recommendations | Full day's safety report                                                          |
| <code>weekly_report()</code>     | Summary, Trends, Frequent Violations, Risk Areas, Recommendations     | Week-over-week trend report                                                       |
| <code>incident_report()</code>   | Summary, Facts, Risks, Immediate Actions, Recommendations             | Single-incident factual writeup, explicitly told not to speculate or assign blame |
| <code>executive_summary()</code> | Overall Status, Critical Findings, Recommendations                    | ≤200-word leadership brief                                                        |

Every prompt explicitly instructs the model to use only the supplied data and not invent or infer information — a sound guardrail for a safety-compliance document. All four return fixed Markdown section headers, which keeps downstream rendering (e.g., a Jinja2 template or the demo UI) predictable.

Observation: each prompt currently interpolates data via `{data}`, i.e. Python's default `str()` of the Pydantic object. This works but produces a verbose `repr()`. Section 8 recommends switching to `data.model_dump_json(indent=2)` for a cleaner, more token-efficient prompt once the schema is corrected.

### 3.3 `models.py` (current / outdated)

Defines three Pydantic models: `Event`, `Statistics`, and `RuleEngineData` (time, events, statistics, `total_count`). This schema describes a generic rule-engine event feed, but it does not match the JSON structure actually produced by the video tracking pipeline (see Section 6 for the full comparison). As delivered, feeding the pipeline's real output into `RuleEngineData.model_validate()` would fail validation.

### 3.4 `llm_main.py` (entry point)

This file is empty. Nothing in the package currently: reads a tracking JSON file from disk, constructs the Pydantic model from it, calls `HSEPromptGenerator`, loads `LLMModel`, or writes/prints the final report. All three other files are library code — none of them are executed automatically. This is the single most important gap for the Demo Team and is addressed directly in Section 8 with a ready-to-use entry-point script.

## 4. Intended Data Flow

Based on the code, the module is designed to operate as follows:

The upstream video-tracking / rule-engine pipeline processes a video and writes a structured JSON file (`video_info`, `tracking_summary`, `alerts_report`).

That JSON is loaded and validated into a Pydantic model.

The validated object is passed into one of the four `HSEPromptGenerator` methods, producing a prompt string.

`download_model()` is run once to cache the LLM locally; `LLMModel` loads it from that local path onto the GPU.

`LLMModel.generate(prompt)` returns the Markdown report text.

The Markdown is stored, displayed, or rendered to PDF/HTML for the end user.

Currently, step 1 produces JSON that does not match the model used in step 2 (Section 6), and steps 2–6 are not wired together anywhere in the code (Section 3.4).

## 5. External Dependencies

Inferred from the import statements (no `requirements.txt` was included in the archive):

- `pydantic`
- `torch` (with CUDA support)

transformers •

huggingface\_hub •

A GPU with sufficient VRAM for Qwen2.5-1.5B-Instruct in bfloat16 (roughly 3–4 GB for weights alone, more with activations/context) is required; there is no quantized or CPU path in the current code.

## 6. Schema Reconciliation — Current vs. Target JSON

The Demo Team supplied the following target JSON as the actual output of the video-tracking pipeline:

```
{
  "video_info": {
    "source": "warehouse.mp4",
    "fps": 30,
    "total_frames": 900,
    "processed_at": "2026-08-03T12:00:00"
  },
  "tracking_summary": {
    "object_durations": { "12": {"class": "Person", "duration_seconds": 24.3} },
    "object_movement": { "12": {"total_distance": 812.4, "avg_speed": 3.1} },
    "classes": { "0": "Hardhat", "5": "Person", "9": "vehicle" }
  },
  "alerts_report": {
    "total_alerts": 12,
    "alerts_by_type": {"machinery_person_proximity": 3, "person_no_hardhat": 9},
    "severity_distribution": {"critical": 3, "medium": 9},
    "alerts": [
      {
        "type": "machinery_person_proximity",
        "frame_id": 145,
        "track_id": 12,
        "machinery_id": 7,
        "distance_m": 0.62,
        "severity": "critical",
        "message": "Machinery 7 approached person 12! Distance: 0.62m"
      }
    ]
  }
}
```

This has a fundamentally different shape from models.py's RuleEngineData. The mismatch:

| models.py (current) | Target JSON                                                            | Verdict                                                  |
|---------------------|------------------------------------------------------------------------|----------------------------------------------------------|
| RuleEngineData.time | video_info.processed_at /<br>video_info.source / fps /<br>total_frames | No equivalent — replaced by a dedicated VideoInfo object |

| models.py (current)                                                            | Target JSON                                                                                                                                        | Verdict                                                                                                                                       |
|--------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------|
| RuleEngineData.events:<br>List[Event]                                          | alerts_report.alerts: List[Alert]                                                                                                                  | Renamed and restructured — Alert has type/frame_id/track_id/severity/message plus optional machinery_id/distance_m, not Event's zone/duration |
| RuleEngineData.statistics:<br>Statistics (persons, forklifts, pallets, alerts) | tracking_summary<br>(object_durations,<br>object_movement, classes) +<br>alerts_report (total_alerts,<br>alerts_by_type,<br>severity_distribution) | Fixed per-class counters replaced by dynamic, ID-keyed dictionaries — cannot reuse Statistics as-is                                           |
| RuleEngineData.total_count                                                     | alerts_report.total_alerts                                                                                                                         | Renamed, and scoped specifically to alerts rather than the whole document                                                                     |

Conclusion: models.py must be replaced, not patched. Section 7 provides the corrected model set (delivered separately as models\_v2.py), which validates the target JSON exactly, including its dynamically-keyed dictionaries (object IDs and class IDs as string keys) and the optional machinery\_id / distance\_m fields that only apply to proximity-type alerts.

## 7. Corrected Pydantic Schema (models\_v2.py)

The following model set is provided as a separate file, models\_v2.py, ready to replace models.py. It was structurally checked field-by-field against the sample JSON above (pydantic could not be executed in this offline review environment, but every key and type in the sample was manually matched against the model definitions below).

```

from datetime import datetime
from typing import Dict, List, Optional
from pydantic import BaseModel, Field

class VideoInfo(BaseModel):
    source: str
    fps: int
    total_frames: int
    processed_at: datetime

class ObjectDuration(BaseModel):
    class_: str = Field(alias="class")
    duration_seconds: float
    class Config:
        populate_by_name = True

class ObjectMovement(BaseModel):

```

```

total_distance: float
avg_speed: float

class TrackingSummary(BaseModel):
    object_durations: Dict[str, ObjectDuration]
    object_movement: Dict[str, ObjectMovement]
    classes: Dict[str, str]

class Alert(BaseModel):
    type: str
    frame_id: int
    track_id: int
    severity: str
    message: str
    machinery_id: Optional[int] = None
    distance_m: Optional[float] = None

class AlertsReport(BaseModel):
    total_alerts: int
    alerts_by_type: Dict[str, int]
    severity_distribution: Dict[str, int]
    alerts: List[Alert]

class PipelineReport(BaseModel):
    video_info: VideoInfo
    tracking_summary: TrackingSummary
    alerts_report: AlertsReport

```

Design notes:

object\_durations, object\_movement, and classes use Dict[str, ...] because their keys are dynamic track/class IDs, not fixed field names — this is what actually lets arbitrary object counts (12, 7, or 200 tracked objects) validate without changing the model.

machinery\_id and distance\_m are Optional on Alert because they only appear on machinery\_person\_proximity alerts; a person\_no\_hardhat alert will validate fine without them.

class\_ uses a Pydantic alias ("class") because class is a reserved Python keyword and cannot be a field name directly.

processed\_at is typed as datetime rather than str, so Pydantic parses and validates the ISO-8601 timestamp automatically.

prompt.py's four method signatures (data: RuleEngineData) should be updated to data: PipelineReport once models.py is replaced; no other change to prompt.py is required, since the templates simply interpolate whatever data object they're given.

## 8. How the Demo Team Should Use This Module

This section is a self-contained runbook to get from a tracking-pipeline JSON file to a rendered Markdown safety report.

## 8.1 One-time environment setup

Use a machine/VM with an NVIDIA GPU and a CUDA-enabled PyTorch build (LLMModel raises `RuntimeError` immediately without one).

Install dependencies: `pip install pydantic torch transformers huggingface_hub`

Replace `models.py` with the corrected `models_v2.py` from Section 7 (or rename it to `models.py` in the project).

Update `prompt.py`'s four type hints from `RuleEngineData` to `PipelineReport` (a find-and-replace).

## 8.2 One-time model download

```
from llm_client import download_model

download_model(path="LLM/model", model_name="Qwen/Qwen2.5-1.5B-Instruct")
# Run once; caches the model locally so future runs don't re-download it.
```

## 8.3 Minimal entry-point script (fills the empty `llm_main.py`)

Since `llm_main.py` ships empty, the Demo Team needs a script like this to actually run the pipeline end-to-end:

```
import json
from llm_client import LLMModel
from models_v2 import PipelineReport
from prompt import HSEPromptGenerator

# 1. Load the tracking pipeline's JSON output
with open("tracking_output.json") as f:
    raw = json.load(f)

# 2. Validate it into the corrected schema
report_data = PipelineReport.model_validate(raw)

# 3. Build the prompt for the report type you need
prompt = HSEPromptGenerator.daily_report(report_data)
# Other options: weekly_report / incident_report / executive_summary

# 4. Load the model (one-time cost per process) and generate
llm = LLMModel(model_path="LLM/model")
markdown_report = llm.generate(prompt, max_new_tokens=600)

# 5. Save or display the result
with open("daily_report.md", "w") as f:
```

```
f.write(markdown_report)
print(markdown_report)
```

## 8.4 Choosing a report type for the demo

| Demo scenario                                 | Call this method                                               |
|-----------------------------------------------|----------------------------------------------------------------|
| Full end-of-day walkthrough                   | <code>HSEPromptGenerator.daily_report(report_data)</code>      |
| Trend/pattern view across multiple days' data | <code>HSEPromptGenerator.weekly_report(report_data)</code>     |
| Drilling into one specific alert              | <code>HSEPromptGenerator.incident_report(report_data)</code>   |
| Short leadership-facing recap                 | <code>HSEPromptGenerator.executive_summary(report_data)</code> |

## 8.5 What to check before/during the demo

Confirm GPU + CUDA are available on the demo machine ahead of time — there is no CPU fallback, so this will fail live if skipped. •

Pre-run the model download step before the demo session; `snapshot_download` over a conference-room network can be slow or fail. •

Validate the tracking JSON with `PipelineReport.model_validate()` before the demo so schema errors surface early, not on stage. •

Since generation is deterministic (`do_sample=False`), the same input JSON will always produce the same report — useful for rehearsing the exact demo output in advance. •

`max_new_tokens` defaults to 300 in `generate()`; the daily/weekly templates ask for five sections and will likely need 500–800 tokens to avoid truncation — raise the value explicitly, as shown in 8.3. •

## 9. Recommendations

Ship a real `llm_main.py` (Section 8.3) so the module is runnable as a package, not just a set of importable components. ١١

Replace `models.py` with `models_v2.py` and update `prompt.py`'s four type hints accordingly (Section 7). ١٢

Add a `requirements.txt` (`pydantic`, `torch`, `transformers`, `huggingface_hub`) so environment setup is reproducible. ١٣

Switch the prompt templates' `{data}` interpolation to `data.model_dump_json(indent=2)` for a cleaner, cheaper prompt once `PipelineReport` is in place. ١٤

Add basic error handling around model loading and JSON validation (currently a malformed tracking JSON or a missing GPU will raise a raw exception with no operator-facing message). ١٥

Consider a documented CPU or quantized fallback path if the module may ever need to run on non-GPU demo hardware. 19

## 10. Conclusion

The model-loading and prompt-generation logic in `llm_client.py` and `prompt.py` is solid, safety-conscious (deterministic decoding, explicit no-speculation instructions), and ready to use. The two blocking gaps — an empty entry point and an outdated schema — are both addressed directly in this report: Section 8.3 supplies a working `llm_main.py`, and Section 7 supplies `models_v2.py`, a corrected Pydantic schema that validates the pipeline's actual `video_info` / `tracking_summary` / `alerts_report` JSON exactly, including its dynamic per-object and per-class dictionaries. With those two pieces in place, the Demo Team can run the four-step flow in Section 8.1–8.3 end-to-end on real tracking output.